

Classic ARIMA Model - Baseline

We used ARIMA as our baseline to forecast future stock prices.

```
In [19]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.stattools import adfuller
from statsmodels.graphics.tsaplots import plot_pacf, plot_acf
from statsmodels.tsa.arima.model import ARIMA
from sklearn.ensemble import RandomForestRegressor
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error
```

Load and Preprocess Data

```
In [20]: # Load and preprocess data
df = pd.read_csv('./berkshire_hathaway_data.csv', index_col=0, parse_dates=True)
df = df[df.index >= '2020-01-01']
df = df.asfreq('D').fillna(method='ffill')
df['log_Close'] = np.log(df['Close'])
df['date_ordinal'] = df.index.map(lambda date: date.toordinal())
print(df)
```

	Open	High	Low	Close	Adj Close	\
Date						
2020-01-02	227.509995	228.389999	226.710007	228.389999	228.389999	
2020-01-03	225.690002	227.429993	225.479996	226.179993	226.179993	
2020-01-04	225.690002	227.429993	225.479996	226.179993	226.179993	
2020-01-05	225.690002	227.429993	225.479996	226.179993	226.179993	
2020-01-06	224.990005	227.130005	224.699997	226.990005	226.990005	
...	
2024-07-25	431.600006	439.630005	431.600006	433.290009	433.290009	
2024-07-26	435.660004	439.000000	434.100006	437.660004	437.660004	
2024-07-27	435.660004	439.000000	434.100006	437.660004	437.660004	
2024-07-28	435.660004	439.000000	434.100006	437.660004	437.660004	
2024-07-29	438.010010	439.440002	435.489990	438.309998	438.309998	

	Volume	log_Close	date_ordinal
Date			
2020-01-02	3764000.0	5.431055	737426
2020-01-03	3023900.0	5.421331	737427
2020-01-04	3023900.0	5.421331	737428
2020-01-05	3023900.0	5.421331	737429
2020-01-06	4263000.0	5.424906	737430
...
2024-07-25	3334300.0	6.071407	739092
2024-07-26	2717600.0	6.081442	739093
2024-07-27	2717600.0	6.081442	739094
2024-07-28	2717600.0	6.081442	739095
2024-07-29	2620000.0	6.082926	739096

[1671 rows x 8 columns]

Stationarity Handler

To make sure ARIMA works properly, we remove trends and make the data more stationary.

```
In [21]: result = adfuller(df['log_Close'].dropna())

print("Test Statistic:", result[0])
print("P-Value:", result[1])

df['diff_log_Close'] = df['log_Close'].diff().dropna()
if(df['diff_log_Close'][1] < 0.05):
    print("Stationary | D=1")
else:
    print("Non-Stationary")

print(result)

# Plot
df['diff_log_Close'].plot(figsize=(15,8))
```


Test Statistic: -0.5218750340147723

P-Value: 0.8876765782218701

Stationary | D=1

(-0.5218750340147723, 0.8876765782218701, 11, 1659, {'1%': -3.4342978282123258, '5%': -2.8632837412222885, '10%': -2.567698326213784}, -10057.536085265187)

Out[21]: <Axes: xlabel='Date'>


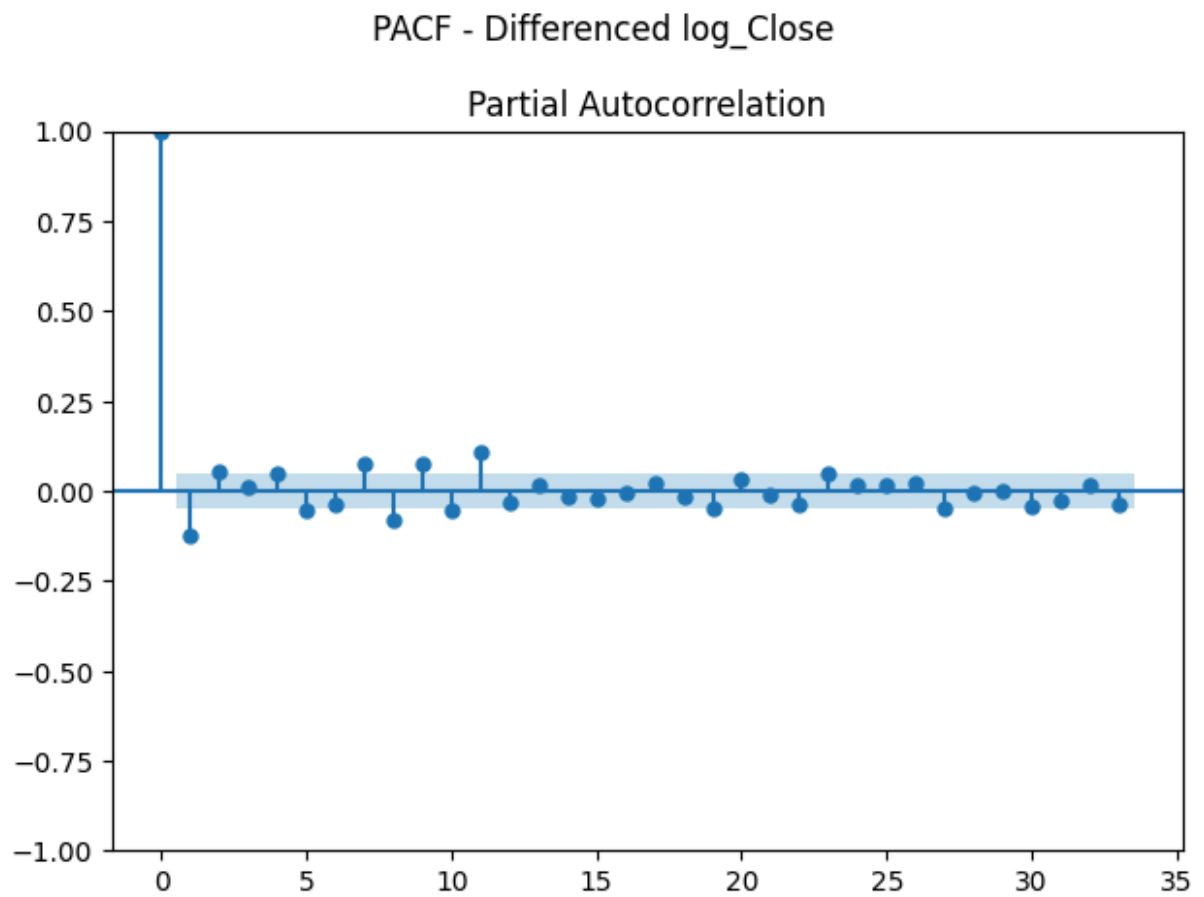


Plot PACF + ACF

The purpose of this is to

```
In [22]: # PACF
fig_pacf = plot_pacf(df['diff_log_Close'].dropna(), method='ywm')
fig_pacf.suptitle("PACF - Differenced log_Close")
fig_pacf.tight_layout()
plt.show()

# ACF
fig_acf = plot_acf(df['diff_log_Close'].dropna())
fig_acf.suptitle("ACF - Differenced log_Close")
fig_acf.tight_layout()
plt.show()
```



ARIMA Train Model

Using the information provided from the graphs, we are going to use this to train our model. For our testing, we decided to predict using the last 30 days to predict the next 30 days.

```
In [23]: train = df.iloc[:-30]
test = df.iloc[-30:]

train_idx = df.index <= train.index[-1]
test_idx = df.index > train.index[-1]

arima = ARIMA(train['Close'], order=(0, 1, 0)).fit()
fig, ax = plt.subplots(figsize=(15,8))

# Plot the real values of stock prices
ax.plot(df['Close'], label='Actual Return')

# Plot fitted values of our model
train_pred = arima.fittedvalues
ax.plot(train.index, train_pred, color='green', label='fitted')

prediction_result = arima.get_forecast(30)
conf_int = prediction_result.conf_int()

# Plot lower and upper limits
lower = conf_int[conf_int.columns[0]]
upper = conf_int[conf_int.columns[1]]

forecast = prediction_result.predicted_mean

ax.plot(test.index, forecast, label='forecast')
ax.fill_between(test.index, lower, upper, color='red', alpha=0.3)
ax.legend()
```

```
Out[23]: <matplotlib.legend.Legend at 0x2111892d650>
```



Evaluate Model

Using RMSE and MAE metrics to see the performance of our model

```
In [24]: y_true = test['Close'].values
rmse = np.sqrt(mean_squared_error(y_true, forecast))
mae = mean_absolute_error(y_true, forecast)

print(f'RMSE: {rmse}')
print(f'MAE: {mae}')
```

RMSE: 22.248062804324675
MAE: 18.150012207031214

Linear Regression

```
In [25]: lr_model = LinearRegression()
lr_model.fit(train[['date_ordinal']], train['Close'])
lr_forecast = lr_model.predict(test[['date_ordinal']])
```

Compute RMSE and MAE

```
In [26]: lr_rmse = np.sqrt(mean_squared_error(y_true, lr_forecast))
lr_mae = mean_absolute_error(y_true, lr_forecast)
print(f'RMSE: {lr_rmse}')
print(f'MAE: {lr_mae}')
```

RMSE: 31.580726027481326
MAE: 29.103639338446875

Plot Linear Regression

```
In [27]: plt.figure(figsize=(15,10))
plt.plot(df.index, df['Close'], label='Actual Close', alpha=0.6)
plt.plot(train.index, train_pred, label='Train Fit (LR)', color='green')
plt.plot(test.index, lr_forecast, label='Test Forecast (LR)', color='blue')
plt.title('Linear Regression Model - Close Price Prediction')
plt.legend()
plt.grid(True)
```



Random Forest

```
In [28]: rf_model = RandomForestRegressor(n_estimators=100, random_state=42)
rf_model.fit(train[['date_ordinal']], train['Close'])
rf_forecast = rf_model.predict(test[['date_ordinal']])

print(rf_forecast)
```

```
[406.98698975 406.98698975 406.98698975 406.98698975 406.98698975
406.98698975 406.98698975 406.98698975 406.98698975 406.98698975
406.98698975 406.98698975 406.98698975 406.98698975 406.98698975
406.98698975 406.98698975 406.98698975 406.98698975 406.98698975]
```

Compute RMSE + MAE


```
In [29]: rf_rmse = np.sqrt(mean_squared_error(y_true, rf_forecast))
rf_mae = mean_absolute_error(y_true, rf_forecast)
print(f'RMSE: {rf_rmse}')
print(f'MAE: {rf_mae}')
```

RMSE: 22.097841493564125

MAE: 18.012877441406218

Plot Random Forest

```
In [30]: plt.figure(figsize=(15,10))
plt.plot(df.index, df['Close'], label='Actual Close', alpha=0.6)
plt.plot(train.index, train_pred, label='Train Fit (RF)', color='green')
plt.plot(test.index, rf_forecast, label='Test Forecast (RF)', color='blue')
plt.title('Random Forest Model - Close Price Prediction')
plt.legend()
plt.grid(True)
```



Compare Baseline Models

```
In [31]: print("ARIMA")
print(f'RMSE: {rmse}')
print(f'MAE: {mae}\n')

print("Linear Regression")
print(f'RMSE: {lr_rmse}')
print(f'MAE: {lr_mae}\n')
```



```
print("Random Forest")
print(f'RMSE: {rf_rmse}')
print(f'MAE: {rf_mae}')
```

ARIMA

RMSE: 22.248062804324675

MAE: 18.150012207031214

Linear Regression

RMSE: 31.580726027481326

MAE: 29.103639338446875

Random Forest

RMSE: 22.097841493564125

MAE: 18.012877441406218

```
In [32]: fig, axs = plt.subplots(3, 2, figsize=(8, 8))

# ARIMA Graph
arima_error = y_true - forecast
axs[0, 0].plot(test.index, arima_error**2, label='ARIMA Squared Error')
axs[0, 0].set_title('ARIMA - Squared Error')
axs[0, 0].legend()

axs[0, 1].plot(test.index, np.abs(arima_error), label='ARIMA Absolute Error', color='red')
axs[0, 1].set_title('ARIMA - Absolute Error')
axs[0, 1].legend()

# Linear Regression Graph
lr_error = y_true - lr_forecast
axs[1, 0].plot(test.index, lr_error**2, label='Linear Regression Squared Error')
axs[1, 0].set_title('Linear Regression - Squared Error')
axs[1, 0].legend()

axs[1, 1].plot(test.index, np.abs(lr_error), label='Linear Regression Absolute Error', color='red')
axs[1, 1].set_title('Linear Regression - Absolute Error')
axs[1, 1].legend()

# Random Forest Graph
rf_error = y_true - rf_forecast
axs[2, 0].plot(test.index, rf_error**2, label='Random Forest Squared Error')
axs[2, 0].set_title('Random Forest - Squared Error')
axs[2, 0].legend()

axs[2, 1].plot(test.index, np.abs(rf_error), label='Random Forest Absolute Error', color='red')
axs[2, 1].set_title('Random Forest - Absolute Error')
axs[2, 1].legend()

plt.tight_layout()
plt.show()
```

